

What Gets a Reply

A pre-registered study of 12,077 outreach emails: what makes a cold email get a reply, how many follow-ups are worth sending — and which of our answers survived a test on sealed data they had never seen.

9

held / confirmed

4

direction held

4

too close to call

2

failed the test

1

untestable

Executive summary

Two findings are confirmed — meaning they were found in 2025, written down as predictions, and then re-appeared in sealed 2026 data at the predicted size:

1. **Mass-produced emails get roughly a quarter of the replies.** Sending the same body to three or more people collapsed reply rates in both years (4.8% vs 9.3% in 2025; 1.8% vs 8.0% in 2026) — measured within the same sender, so it is not "worse reps use templates."
2. **A concrete reason to write *now* roughly doubles replies.** A launch, a funding round, an event, a referral — something checkable that answers "why this email, this week." 7.9% vs 4.6% in 2025; 8.5% vs 4.1% in 2026. It beat every other quality we measured, both years, and held on the stricter "showed real interest" outcome.

The most repeated advice did not survive scrutiny. Researched personalisation — citing facts about the recipient, writing that could only fit them — never produced a significant effect in either year. Links don't hurt. Greeting by name does nothing. The "keep it short" rule is a floor at ~100 words, not a virtue of ever-shorter emails.

Two 2025 findings failed the sealed test and are withdrawn — including our most promising one ("recipient-focused emails win", +5.3pp in 2025, -0.4pp in 2026). That is the safeguard working: without the sealed year, both would be in the recommendations.

Follow-ups: the first follow-up is worth as much as the opener; ~90% of all eventual replies arrive by the third touch. Both patterns repeated in 2026.

What the data supports doing — graded by evidence:

- **Strong:** stop sending identical bodies at scale from personal mailboxes; require every cold email to name its reason-to-write-now.
- **Moderate:** keep openers under ~100 words; always send follow-ups 1–2, usually 3; stop after 4 unless there's a new reason.
- **Time back:** stop polishing links out, stop first-name rituals, stop research-heavy personalisation as a reply tactic — none of it moved replies.

The question, and why naive answers mislead

The question sounds simple: *what makes someone reply to a cold email?* Getting a defensible answer is not, for three reasons.

1.1 • Good reps poison simple comparisons

Reps choose who gets which email. If strong reps write short emails to warm prospects, "short emails win" may just mean "good reps win." Our primary comparison is therefore **within-sender**: each rep's emails with a trait against the same rep's emails without it. Rep skill and rep targeting drop out of the comparison.

1.2 • Test enough traits and some "win" by luck

We measured ~40 traits; test that many and 1–2 "win" by chance. Every result is corrected for this, within pre-declared families of tests. We also re-ran the whole analysis on deliberately scrambled data ten times: it correctly found almost nothing (5 spurious hits in 240 tests), so the machinery itself doesn't manufacture findings.

1.3 • Staring at data until a pattern appears guarantees a pattern

The deepest problem: explore data freely and you *will* find patterns — some real, some noise you fit to. The only clean defence is the one machine-learning uses: **a training set and a validation set**. We explored 2025 freely, wrote every conclusion down as a testable prediction with its pass/fail rule, committed them to a timestamped record — and only then opened Jan–Jul 2026, once. No tuning, no second look.

How to read "failed" and "held." Our split is by time, not random, so 2026 is a genuinely different world (different sender mix, different tooling). A finding that *fails* has two possible readings — it was never real, or it was real and the world changed — and we deliberately do not rescue any failure with the second reading: failed means *don't build on it*. A finding that *holds* cleared a harder bar than a random split: it survived fresh noise *and* a shifted world. That is the property you want before acting on it, because every decision applies to the future — also a changed world.


How the study worked

From ~117,000 sent-email records: 12,077 qualifying "first touches" (fresh outreach to a person after ≥30 days of silence, Gmail-recorded, from commercial senders; bounces and pre-existing threads excluded). Every reply candidate was classified by AI (human vs auto-reply vs calendar bot vs bounce — validated at 99% agreement); only genuine human replies count. Two outcomes per email: **replied** (a human wrote back) and **interested** (the reply shows forward motion — validated at 92.8%).

Layer 1 — countable traits (length, questions, links, subject shape, name use, template detection) measured by audited code. **Layer 2 — judged qualities**: an AI read each of 12,462 emails blind — sender, date, recipient and outcome all redacted — and scored 12 pre-registered qualities against a rubric with written anchors. A second scoring pass and an independent second AI measured which qualities can be scored consistently at all (Section 6); two couldn't and were discarded.

The two findings that survived everything


3.1 · Mass-produced emails get a quarter of the replies



Reply rate, cold pitches. "Templated" = same body sent to 3+ people. Within-sender; 2026 $p=0.005$ (replied), $p=0.001$ (interested: 0.9% vs 5.6%). Agreed in 10 of 12 reps (2025), 3 of 4 (2026).

The effect is not subtle, and its fingerprints show up elsewhere: bold text (−4.3pp) and re-using the recipient's first name mid-email (−5.4pp) both looked harmful in 2025, but 77–80% of the emails carrying them were templates — they are substantially the same finding in different clothes, and we count them once, not three times.

3.2 · A real reason to write now roughly doubles replies



Reply rate, cold pitches. "Reason" = an explicit, checkable occasion: launch, funding, event, visit, referral, news. 2026: $p=0.047$ replied / $p=0.012$ interested; effect landed inside the pre-declared band (+1.6 to +7.6); top-ranked judged quality both years, as predicted.

This is the study's sharpest lesson because of what it beat: **research-heavy personalisation did nothing** (next section), while having an *occasion* doubled replies. "I know things about you" and "there is a reason I'm writing to you today" feel similar — the data says only the second one matters. It also survives a length control: it is not "short emails in disguise."


Everything that moved replies in 2025 — and its 2026 verdict

Effect in percentage points on reply rate (dot), statistical uncertainty (bar), within-sender, cold pitches, n=5,153. Verdicts from the sealed 2026 test (n=1,286).

Names a real reason to write now		+4.6	Confirmed
Recipient-focused writing		+5.3	Failed 2026
Under 100 words		+4.5	Direction held
Subject line is a question		+4.5	Direction held
Asks 2+ questions		-2.5	Direction held
Has bold text		-4.3	Mostly templating
Reuses their name mid-email		-5.4	Untestable 2026
Mass template (3+ identical)		-5.4	Confirmed

What does nothing — tested with the same rigor

BELIEF	WHAT WE FOUND
Links hurt deliverability/replies	No. Direction was positive both years (+4.9pp, then +1.5pp); never solid enough to claim the reverse.
Greet them by name	Nothing either year (-2.1pp, then +2.1pp — noise around zero).
Mention their company	Nothing (0.0pp, then -0.4pp). Name-dropping ≠ relevance.
Cite researched facts about them	Nothing significant either year (+1.3, then +0.6).
Make it read written-just-for-them	Nothing significant either year (+1.9, then +2.1).
Nail one crisp, specific ask	Nothing (-0.0, then +0.9).
Shorter is always better	No — it's a floor at ~100 words , not a slope:




These nulls are findings, not absences: each had the sample size to show an effect the size of our confirmed ones, and didn't. They contradict most outbound playbooks, which is precisely what makes them valuable — they say where effort is being wasted.

SECTION 5

Follow-ups: the first one is worth as much as the opener

Chance of a reply after each touch, among people still silent at that point. Vertical bars = uncertainty; few people receive late touches in 2026, so those bars are wide (clipped at 8%). All email types, 8,591 pushes (2025) / 3,486 (2026).




61% of all replies arrive *after* the first email — an unsent follow-up forfeits more than half the outcome. The first follow-up out-performs the opener itself (5.0% vs 4.1% in 2025; 6.0% vs 5.8% in 2026 — predicted and held).

~90% of eventual replies are in by touch 3; touches 5–8 added 0.6 points *combined* in 2025. One honest caveat: who receives touch 5 is the rep's choice, so late-touch rates describe the prospects reps chose to keep chasing.

SECTION 5B

The mirror: what our 12,462 emails actually look like

Blind AI ratings, before any outcome was attached.



The team writes clean, well-formed emails with clear asks — that are mostly about us. The two traits that matter most (a reason-to-write; not being a template) are exactly where the corpus is weakest: 20% of emails have no why-now, and 44% are visible mail-merge.

Two findings died on the sealed data — and were withdrawn

"Recipient-focused emails win" was 2025's most attractive result: +5.3pp, significant after correction, stronger in every robustness check we ran inside 2025. On the sealed 2026 data: **−0.4pp**, with the 2025 effect size statistically excluded. It had been labelled exploratory, moderate-reliability, small-sample — every flag earned. It is out.

"Bullet points ruin event invites" (−9.6pp in 2025) returned +1.3pp in 2026. A one-year artefact, now on record as such.

Failure has two possible readings — never real, or real-then-the-world-changed. We do not use the second as a rescue: failed means withdrawn. But it is why "held" is the stronger claim: survivors cleared fresh noise *and* a changed world.

Can you trust an AI judge? We measured it

Each quality was scored twice by the same AI on 1,246 emails, and independently by a second, stronger AI on 1,000. The table shows agreement *beyond chance* (0 = coin flip, 1 = perfect) on the exact cut used in analysis. Two qualities scored near-chance — by *both* AIs, which agreed with each other exactly as much as one agreed with itself (gap −0.02). That means the questions themselves are ambiguous, not the model — so both were discarded, and nothing in this document rests on them.

JUDGED QUALITY	SELF-AGREEMENT	SECOND AI
Reason to write now	0.72	0.72
Research on the recipient	0.75	0.73
Guessed the reader's problem	0.69	0.64
Specific value claim	0.67	0.62
Customer proof	0.65	0.66
Written just for them	0.56	0.60
Clear ask	0.42	0.57
Polish	0.43	0.46
About the recipient	0.43	0.45
Peer tone — discarded	0.32	0.32
Economy — discarded	0.22	0.26

Above 0.6 = solid · 0.4–0.6 = usable with care · below 0.4 = discarded. Separately, the 12 qualities were checked for being one "halo" score in disguise — they are not (mean cross-correlation 0.19; threshold 0.60), with one exception: research-signal and bespokeness overlap so strongly ($r=0.83$) that they are reported as one finding.

SECTION 7

The full scorecard: 19 predictions, sealed-data verdicts

Every prediction and its pass/fail rule was committed to a timestamped record before 2026 was opened; power-limited rules were tightened (blind) in a pre-registered audit.

Mass-templated emails get fewer replies	Confirmed — both years, both outcomes (p=0.005 / 0.001)
A real reason-to-write-now wins	Confirmed — both outcomes (p=0.047 / 0.012)
Its effect size lands in the pre-declared band	Held (+2.9, band +1.6 to +7.6)
It is the #1 judged quality	Held — ranked first again
Follow-up #1 is worth as much as the opener	Held (5.8% vs 6.0%)
Most replies arrive by the 4th touch	Held (96%)
Same pattern on the “interested” outcome	Held
“Pain”/“proof” negatives are just email length	Held
The two unreliable qualities stay silent	Consistent
Under 100 words beats longer	Direction held; 2025's size (+4.5) excluded by 2026
Question subject lines help	Direction held; too rare in 2026 (n=79) to be sure
Asking 2+ questions hurts	Direction held
Templated event invites do worse	Direction held
Short event invites do better	Direction held; 2025's +17.5 excluded
Bold text hurts	Too close to call (cell shrank; wrong sign inside noise)
Personalisation & crisp asks do nothing	Still nothing — 2026 too noisy to certify $\leq 3pp$
Recipient-focused emails win	Failed — +5.3 became -0.4; 2025 size excluded
Bullet points ruin event invites	Failed — -9.6 became +1.3
Reusing their name mid-email hurts	Habit vanished (2 emails in 2026)

Scope and limitations — what this study cannot say

1. **It covers Gmail-recorded outbound only** — hand-sent and Apollo-sent. Amplemarket sends leave no reliable reply trail and are excluded; they grew from 20% of outreach in 2025 to 71% in 2026. These findings therefore describe a shrinking, more deliberate slice of outbound (~75% of it in 2025, under 30% in 2026) and say nothing about whether the Amplemarket machine works.
2. **The sending tool doesn't change what works.** Tool-sent emails reply about 2 points lower across the board, but the offset is flat: across all 29 traits tested, the tool changed the effect of essentially none (2 of 29, exactly the chance rate). It shifts the baseline, not the recipe — and is likely partly a tracking artefact, so no claim is made either way.
3. **This is correlation, credibly narrowed — not an experiment.** Within-sender comparison removes rep-level differences, but a rep may still save their best emails for their best prospects. Treat effect sizes as upper bounds; treat directions as trustworthy. The only way to do better is an A/B test, which this study's two confirmed findings are the natural candidates for.
4. **Some 2026 verdicts are power-limited.** 2026 is a quarter of 2025; four predictions came back "too close to call" — reported as exactly that, never as confirmation.
5. **Known residuals, documented not hidden:** 155 Amplemarket emails inside scope on a technicality (all analyses re-run without them; nothing changes); ~0.1% of replies missed via colleague hand-offs; non-English emails (0.5%) judged with an English rubric.

Method registration, defect log, per-analysis outputs and this document's source are version-controlled in the study repository (rules/RUN2_PREREGISTRATION.md, docs/13–16). Analysis: within-sender fixed effects with cluster-robust uncertainty; Benjamini–Hochberg correction within pre-registered families; placebo-tested. AI judging: blinded batches, automated leak gate (0 leaks across 12,462 items), dual-model validation. Prepared 15 August 2026.